


Concurrency in Go




并发的产生?

CPU 计算 vs I/O:


Operation	time
2.5GHz的CPU一个时钟周期	0.4ns
1Gbps的网络传输2KB数据时间	20μs
SSD随机读取时间	150μs
磁盘寻道+旋转时间	15ms

Concurrency



VS

Parallelism



并发的载体

Process

独立的地址空间，内核态上下文切换

Thread

内核态上下文切换，切换开销小一些

Coroutine

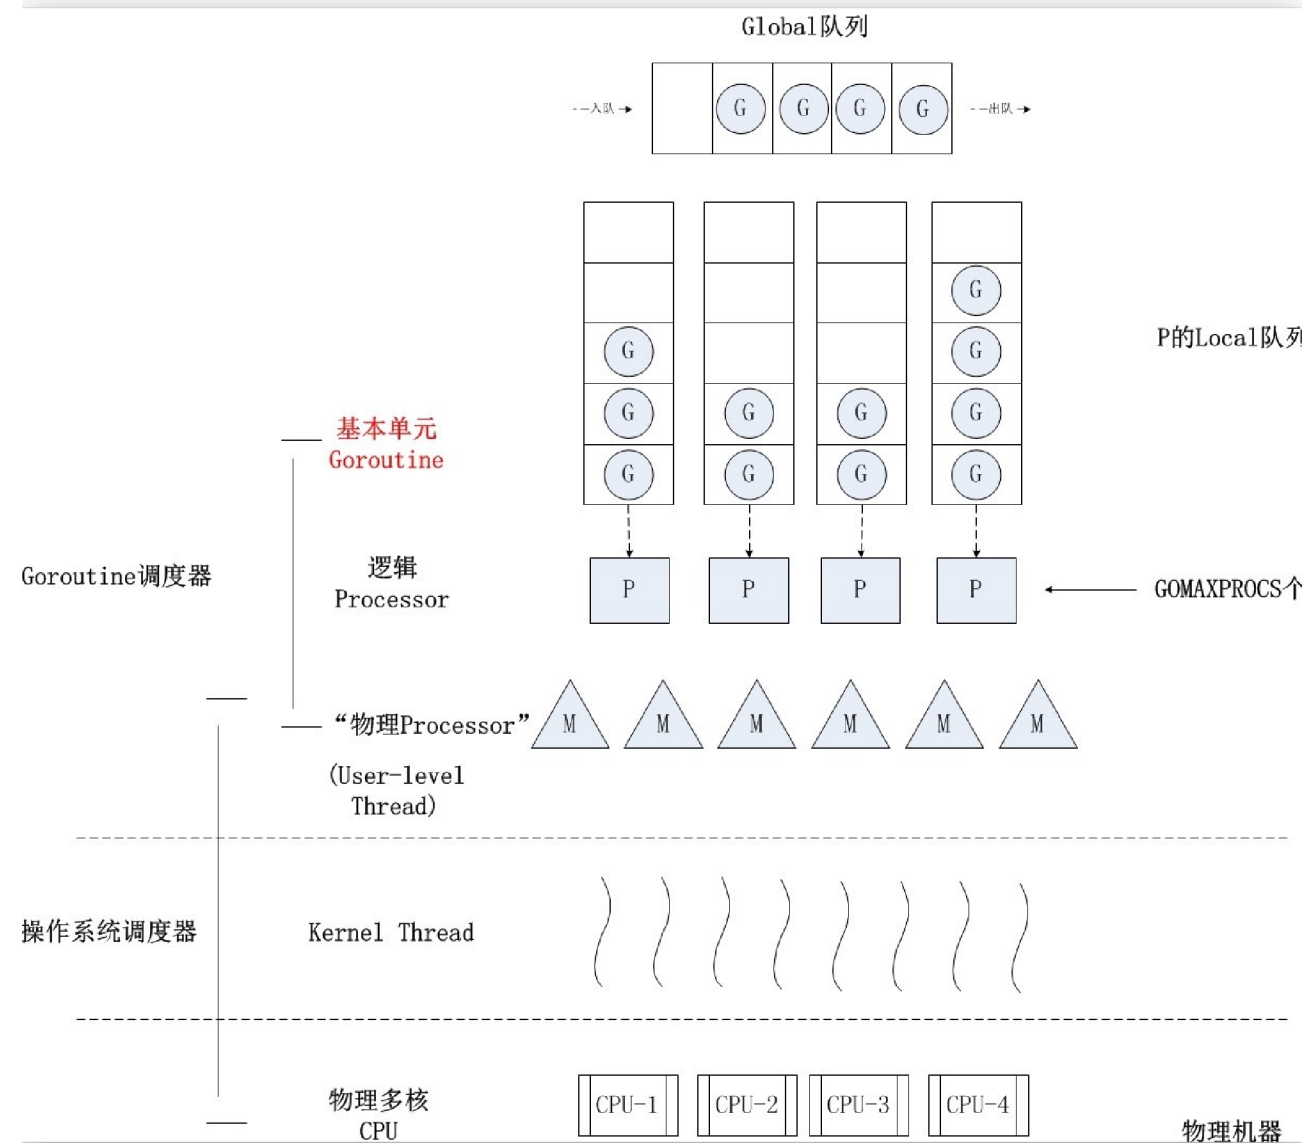
用户态上下文切换

Goroutine

栈轻量、可扩缩容 (2KB ~ 1GB)

由用户态线程调度

Goroutine的调度 - GMP模型



并发从入门到放弃 - Race Condition

计算机中的两个进程同时试图修改一个共享内存的内容，在没有并发控制的情况下，最后的结果依赖于两个进程的执行顺序与时机。而且如果发生了并发访问冲突，则最后的结果是不正确的。(Wikipedia)

并发从入门到放弃 - (1) 丧失原子性

```
var sum int64

func main() {
    for i := 0; i < 1000; i++ {
        go func() {
            sum++
        }()
    }

    fmt.Println(sum)
}
```


并发从入门到放弃 - (2) 丧失顺序性

```
func f1() {  
    fmt.Print(1)  
}  
  
func f2() {  
    fmt.Print(2)  
}  
  
func f3() {  
    fmt.Print(3)  
}  
  
func main() {  
    go f1()  
    go f2()  
    go f3()  
}
```

并发从入门到放弃 - (3) 丧失可见性

```
var s string

func hello() {
    s = "hello"
}

func main() {
    go hello()
    fmt.Print(s)
}
```

原子性、顺序性、可见性都丧失了，还编个JB？

Race检查

Concurrency Primitive

- Sync - Mutex & Atomic
- CSP模型 - Channel

CSP 是 Communicating Sequential Process 的简称，中文可以叫做通信顺序进程，是一种并发编程模型，是一个很强大的并发数据模型，是上个世纪七十年代提出的，用于描述两个独立的并发实体通过共享的通讯 channel(管道)进行通信的并发模型。相对于Actor模型，CSP中channel是第一类对象，它不关注发送消息的实体，而关注与发送消息时使用的channel。

Do not communicate by sharing memory; instead, share memory by communicating. ——Effective Go

拯救刚才的例子 - (1) 原子性

sync/atomic

```
atomic.AddInt64(&sum, 1)
```

sync.Mutex

```
var mx sync.Mutex
```

```
mx.Lock()
```

```
sum++
```

```
mx.Unlock()
```

拯救刚才的例子 - (2) 顺序性

```
func f1() {  
    fmt.Print(1)  
    ch1 <- 1  
}  
  
func f2() {  
    <- ch1  
    fmt.Print(2)  
    ch2 <- 2  
}  
  
func f3() {  
    <- ch2  
    fmt.Print(3)  
    ch3 <- 3  
}
```

拯救刚才的例子 - (3) 可见性

```
func hello() {  
    s = "hello"  
    ch <- 1  
}  
  
func main() {  
    go hello()  
    ch <- ch  
    fmt.Print(s)  
}
```

Understanding Real-World Concurrency Bugs in Go

Application	Shared Memory					Message		Total
	Mutex	atomic	Once	WaitGroup	Cond	chan	Misc.	
Docker	62.62%	1.06%	4.75%	1.70%	0.99%	27.87%	0.99%	1410
Kubernetes	70.34%	1.21%	6.13%	2.68%	0.96%	18.48%	0.20%	3951
etcd	45.01%	0.63%	7.18%	3.95%	0.24%	42.99%	0	2075
CockroachDB	55.90%	0.49%	3.76%	8.57%	1.48%	28.23%	1.57%	3245
gRPC-Go	61.20%	1.15%	4.20%	7.00%	1.65%	23.03%	1.78%	786
BoltDB	70.21%	2.13%	0	0	0	23.40%	4.26%	47

Table 4. Concurrency Primitive Usage. *The Mutex column includes both Mutex and RWMutex.*

Application	Behavior		Cause	
	blocking	non-blocking	shared memory	message passing
Docker	21	23	28	16
Kubernetes	17	17	20	14
etcd	21	16	18	19
CockroachDB	12	16	23	5
gRPC	11	12	12	11
BoltDB	3	2	4	1
Total	85	86	105	66

Table 5. Taxonomy. *This table shows how our studied bugs distribute across different categories and applications.*

Blocking Bugs

Application	Shared Memory			Message Passing		
	Mutex	RWMutex	Wait	Chan	Chan w/	Lib
Docker	9	0	3	5	2	2
Kubernetes	6	2	0	3	6	0
etcd	5	0	0	10	5	1
CockroachDB	4	3	0	5	0	0
gRPC	2	0	0	6	2	1
BoltDB	2	0	0	0	1	0
Total	28	5	3	29	16	4

Table 6. Blocking Bug Causes. *Wait includes both the Wait function in Cond and in WaitGroup. Chan indicates channel operations and Chan w/ means channel operations with other operations. Lib stands for Go libraries related to message passing.*

- Kubernetes

```
1  func finishReq(timeout time.Duration) r ob {
2  -   ch := make(chan ob)
3  +   ch := make(chan ob, 1)
4      go func() {
5          result := fn()
6          ch <- result // block
7      } ()
8      select {
9          case result = <- ch:
10             return result
11          case <- time.After(timeout):
12             return nil
13      }
14 }
```

Figure 1. A blocking bug caused by channel.

- Docker#25384

```
1  var group sync.WaitGroup
2  group.Add(len(pm.plugins))
3  for _, p := range pm.plugins {
4      go func(p *plugin) {
5          defer group.Done()
6      }
7  -   group.Wait()
8  }
9  + group.Wait()
```

Figure 5. A blocking bug caused by WaitGroup.

- context

```
hctx, hcancel := context.WithCancel(ctx)
if timeout > 0 {
    //hcancel.Cancel()
    hctx, hcancel = context.WithTimeout(ctx, timeout)
}
```

```
1 func goroutine1() {  
2     m.Lock()  
3 -   ch <- request //blocks  
4 +   select {  
5 +       case ch <- request  
6 +       default:  
7 +   }  
8     m.Unlock()  
9 }
```

(a) goroutine 1

```
1 func goroutine2() {  
2     for {  
3         m.Lock()    //blocks  
4         m.Unlock()  
5         request <- ch  
6     }  
7 }
```

(b) goroutine 2

Figure 7. A blocking bug caused by wrong usage of channel with lock.

Non-Blocking Bugs

Application	Shared Memory				Message Passing	
	traditional	anon.	waitgroup	lib	chan	lib
Docker	9	6	0	1	6	1
Kubernetes	8	3	1	0	5	0
etcd	9	0	2	2	3	0
CockroachDB	10	1	3	2	0	0
gRPC	8	1	0	1	2	0
BoltDB	2	0	0	0	0	0
Total	46	11	6	6	16	1

Table 9. Root causes of non-blocking bugs. *traditional: traditional non-blocking bugs; anonymous function: non-blocking bugs caused by anonymous function; waitgroup: misusing WaitGroup; lib: Go library; chan: misusing channel.*

- Docker

```
1    for i := 17; i <= 21; i++ { // write
2 -    go func() { /* Create a new goroutine */
3 +    go func(i int) {
4        apiVersion := fmt.Sprintf("v1.%d", i) // read
5        ...
6 -    }()
7 +    }(i)
8    }
```

Figure 8. A data race caused by anonymous function.

- etcd

```
1 func (p *peer) send() {
2     p.mu.Lock()
3     defer p.mu.Unlock()
4     switch p.status {
5         case idle:
6 +         p.wg.Add(1)
7         go func() {
8 -         p.wg.Add(1)
9             ...
10            p.wg.Done()
11        }()
12    case stopped:
13    }
14 }
```

(a) func1

```
1 func (p * peer) stop() {
2     p.mu.Lock()
3     p.status = stopped
4     p.mu.Unlock()
5     p.wg.Wait()
6 }
```

(b) func2

Figure 9. A non-blocking bug caused by misusing WaitGroup.

- Docker

```
1 - select {  
2 -     case <- c.closed:  
3 -     default:  
4 +         Once.Do(func() {  
5             close(c.closed)  
6 +         })  
7 - }
```

Figure 10. A bug caused by closing a channel twice.